

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.

1. Binary
2. Multiclass
3. Multi label

```
from sklearn.datasets import make_circles
n_samples = 1000
X,Y= make_circles(n_samples,noise=0.03,random_state=42)
print(X.shape)
print(Y.shape)
```

```
⇒ (1000, 2)
   (1000,)
```

X

```
⇒ array([[ 0.75424625,  0.23148074],
        [-0.75615888,  0.15325888],
        [-0.81539193,  0.17328203],
        ...,
        [-0.13690036, -0.81001183],
        [ 0.67036156, -0.76750154],
        [ 0.28105665,  0.96382443]])
```

Y[:10]

```
⇒ array([1, 1, 1, 1, 0, 1, 1, 1, 1, 0])
```

```
import pandas as pd
circles= pd.DataFrame({'X1':X[:,0], 'X2':X[:,1], 'label':Y})
circles
```



	X1	X2	label
0	0.754246	0.231481	1
1	-0.756159	0.153259	1
2	-0.815392	0.173282	1
3	-0.393731	0.692883	1
4	0.442208	-0.896723	0
...
995	0.244054	0.944125	0
996	-0.978655	-0.272373	0
997	-0.136900	-0.810012	1
998	0.670362	-0.767502	0
999	0.281057	0.963824	0



1000 rows × 3 columns

Next
steps:[Generate code
with](#)


circles

[View recommended
plots](#)[New interactive
sheet](#)

```
import matplotlib.pyplot as plt
plt.scatter(X[:,0],X[:,1],c=Y,cmap=plt.cm.RdYlBu)
```



<matplotlib.collections.PathCollection at 0x7b81e28f8af0>




 Generate

print hello world using rot13

Close 

X.shape,Y.shape

 ((1000, 2), (1000,))

X[0],Y[0]


 (array([0.75424625, 0.23148074]), 1)

```
import tensorflow as tf
tf.random.set_seed(42)
model_1=tf.keras.Sequential([
    tf.keras.layers.Dense(1)
])
model_1.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.SGD(),
                metrics=["accuracy"])
```

model_1.fit(X,Y,epochs=200,verbose=0)

 <keras.src.callbacks.history.History at 0x7b818d09aef0>

model_1.evaluate(X,Y)


 32/32 ————— 0s 1ms/step - accuracy: 0.4955 - loss: 8.1322
[8.059046745300293, 0.5]

```
import tensorflow as tf
tf.random.set_seed(42)
model_2=tf.keras.Sequential([
    tf.keras.layers.Dense(1),
    tf.keras.layers.Dense(1)
])
model_2.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.SGD(),
                metrics=["accuracy"])
```

model_2.fit(X,Y,epochs=200,verbose=0)

 <keras.src.callbacks.history.History at 0x7b818d071180>

model_2.evaluate(X,Y)

 32/32 ————— 0s 2ms/step - accuracy: 0.4955 - loss: 0.6932
[0.6932107210159302, 0.5]

```
import tensorflow as tf
tf.random.set_seed(42)
```

```

model_3=tf.keras.Sequential([
    tf.keras.layers.Dense(100),
    tf.keras.layers.Dense(10),
    tf.keras.layers.Dense(1)

])
model_3.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(),
                metrics=["accuracy"])

```

```
model_3.fit(X,Y,epochs=100,verbose=0)
```

↩ <keras.src.callbacks.history.History at 0x7b818c2ff340>

```
model_3.evaluate(X,Y)
```

↩ 32/32 ————— 0s 2ms/step - accuracy: 0.5036 - loss: 0.6917
[0.6945011019706726, 0.48500001430511475]

1. create a meshgrid for different x values
2. make predictions across the meshgrid
3. plot the predictions as well as line b/w zeros


```

import numpy as np
def plot_decision_boundary(model,X,Y):
    # plt.figure(figsize=(12,8))
    # X=tf.cast(X,tf.float32)
    # Y=tf.cast(Y,tf.float32)
    x_min,x_max=tf.reduce_min(X[:,0]),tf.reduce_max(X[:,0])
    y_min,y_max=tf.reduce_min(X[:,1]),tf.reduce_max(X[:,1])
    xx,yy=np.meshgrid(np.linspace(x_min,x_max,100),np.linspace(y_min,y_max,100))
    x_in = np.c_[xx.ravel(),yy.ravel()]#stack 2d arrays together
    y_pred = model.predict(x_in)
    if(len(y_pred[0])) >1:
        print("doing multiclass classification")
        y_pred=np.argmax(y_pred,axis=1).reshape(xx.shape)
    else:
        print("doing binary classification")
        y_pred=np.round(y_pred).reshape(xx.shape)
    plt.contourf(xx,yy,y_pred.reshape(xx.shape),cmap=plt.cm.RdYlBu,alpha=0.7)
    plt.scatter(X[:,0],X[:,1],c=Y,s=40,cmap=plt.cm.RdYlBu)
    plt.xlim(xx.min(),xx.max())
    plt.ylim(yy.min(),yy.max())

```

```
plot_decision_boundary(model_3,X,Y)
```

313/313 0s 1ms/step
doing binary classification



Suggested code may be subject to a licence |

```
tf.random.set_seed(42)
```

```
model_4=tf.keras.Sequential([
```

```
    tf.keras.layers.Dense(1,activation="linear"),
```

```
])
```

```
model_4.compile(loss=tf.keras.losses.BinaryCrossentropy(),
```

```
                optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
```

```
                metrics=["accuracy"])
```


```
history= model_4.fit(X,Y,epochs=100)
```



```
32/32 ————— 0s 2ms/step - accuracy: 0.5093 - loss: 0.6923
Epoch 82/100
32/32 ————— 0s 1ms/step - accuracy: 0.5093 - loss: 0.6922
Epoch 83/100
32/32 ————— 0s 1ms/step - accuracy: 0.5074 - loss: 0.6922
Epoch 84/100
32/32 ————— 0s 1ms/step - accuracy: 0.5065 - loss: 0.6922
Epoch 85/100
32/32 ————— 0s 1ms/step - accuracy: 0.5065 - loss: 0.6922
Epoch 86/100
32/32 ————— 0s 2ms/step - accuracy: 0.5065 - loss: 0.6922
Epoch 87/100
32/32 ————— 0s 1ms/step - accuracy: 0.5065 - loss: 0.6922
Epoch 88/100
32/32 ————— 0s 1ms/step - accuracy: 0.5065 - loss: 0.6922
Epoch 89/100
32/32 ————— 0s 1ms/step - accuracy: 0.5090 - loss: 0.6922
Epoch 90/100
32/32 ————— 0s 1ms/step - accuracy: 0.5090 - loss: 0.6922
Epoch 91/100
32/32 ————— 0s 1ms/step - accuracy: 0.5090 - loss: 0.6922
Epoch 92/100
32/32 ————— 0s 1ms/step - accuracy: 0.5090 - loss: 0.6922
Epoch 93/100
32/32 ————— 0s 1ms/step - accuracy: 0.5101 - loss: 0.6923
Epoch 94/100
32/32 ————— 0s 2ms/step - accuracy: 0.5103 - loss: 0.6923
Epoch 95/100
32/32 ————— 0s 1ms/step - accuracy: 0.5095 - loss: 0.6923
Epoch 96/100
32/32 ————— 0s 1ms/step - accuracy: 0.5095 - loss: 0.6923
Epoch 97/100
32/32 ————— 0s 2ms/step - accuracy: 0.5095 - loss: 0.6923
Epoch 98/100
32/32 ————— 0s 1ms/step - accuracy: 0.5106 - loss: 0.6923
Epoch 99/100
32/32 ————— 0s 1ms/step - accuracy: 0.5106 - loss: 0.6924
Epoch 100/100
32/32 ————— 0s 1ms/step - accuracy: 0.5126 - loss: 0.6924
```

```
plt.scatter(X[:,0],X[:,1],c=Y,cmap=plt.cm.RdYlBu)
```

 <matplotlib.collections.PathCollection at 0x7b817ff47700>





 Generate


a slider using jupyter widgets



Close

plot_decision_boundary(model_4,X,Y)

 313/313  1s 3ms/step
doing binary classification



```
#non linear activation function
tf.random.set_seed(42)
model_5=tf.keras.Sequential([
    tf.keras.layers.Dense(1,activation="relu"),
])
model_5.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
                metrics=["accuracy"])

model_5.fit(X,Y,epochs=100)
```



Epoch 96/100

32/32 ————— 0s 2ms/step - accuracy: 0.4549 - loss: 4.7938

Epoch 97/100

32/32 ————— 0s 2ms/step - accuracy: 0.4549 - loss: 4.7900

Epoch 98/100

32/32 ————— 0s 1ms/step - accuracy: 0.4549 - loss: 4.7870

Epoch 99/100

32/32 ————— 0s 1ms/step - accuracy: 0.4549 - loss: 4.7844

Epoch 100/100

32/32 ————— 0s 1ms/step - accuracy: 0.4549 - loss: 4.7820

<keras.src.callbacks.history.History at 0x7b818571b6d0>

tf.random.set_seed(42)

model_6=tf.keras.Sequential([

tf.keras.layers.Dense(4,activation='relu'),

tf.keras.layers.Dense(4,activation='relu'),

tf.keras.layers.Dense(1,activation="sigmoid"),

])

model_6.compile(loss=tf.keras.losses.BinaryCrossentropy(),


optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),

metrics=["accuracy"])

model_6.fit(X,Y,epochs=300,verbose=0)

↗ <keras.src.callbacks.history.History at 0x7b8185c252d0>

plot_decision_boundary(model_6,X,Y)

↗ 313/313 ————— 0s 1ms/step
doing binary classification

model_6.evaluate(X,Y)

32/32 — 0s 1ms/step - accuracy: 0.9848 - loss: 0.0903
[0.0886337161064148, 0.984000027179718]

```
#sigmoid
```

```
A= tf.cast(tf.range(-10,10),tf.float32)
```

```
def sigmoid(x):
```


```
    return 1/(1+tf.exp(-x))
```

```
sigmoid(A)
```

```
<tf.Tensor: shape=(20,), dtype=float32, numpy=
array([4.5397872e-05, 1.2339458e-04, 3.3535014e-04, 9.1105117e-04,
       2.4726233e-03, 6.6928510e-03, 1.7986210e-02, 4.7425874e-02,
       1.1920292e-01, 2.6894143e-01, 5.0000000e-01, 7.3105860e-01,
       8.8079703e-01, 9.5257413e-01, 9.8201376e-01, 9.9330717e-01,
       9.9752742e-01, 9.9908900e-01, 9.9966466e-01, 9.9987662e-01],
      dtype=float32)>
```

```
plt.plot(sigmoid(A))
```

```
[<matplotlib.lines.Line2D at 0x7b8185718f40>]
```




```
#relu
```

```
def relu(x):
```

```
    return tf.maximum(0,x)
```

```
plt.plot(relu(A))
```


[<matplotlib.lines.Line2D at 0x7b8184267490>]



```
#tanh
def tanh(x):
    return (tf.exp(x)-tf.exp(-x))/(tf.exp(x)+tf.exp(-x))
```

```
plt.plot(tanh(A))
```


[<matplotlib.lines.Line2D at 0x7b818c2ff310>]



linear activation function returns the tensor unmodified

```
plt.plot(tf.keras.activations.linear(A))
```

↗ [<matplotlib.lines.Line2D at 0x7b8185848be0>]



IMPROVING OUR MODEL

```
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=0)
```

```
tf.random.set_seed(42)
model_7=tf.keras.Sequential([
    tf.keras.layers.Dense(4,activation='relu'),
    tf.keras.layers.Dense(4,activation='relu'),
    tf.keras.layers.Dense(1,activation="sigmoid"),
])
model_7.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
                metrics=["accuracy"])
```


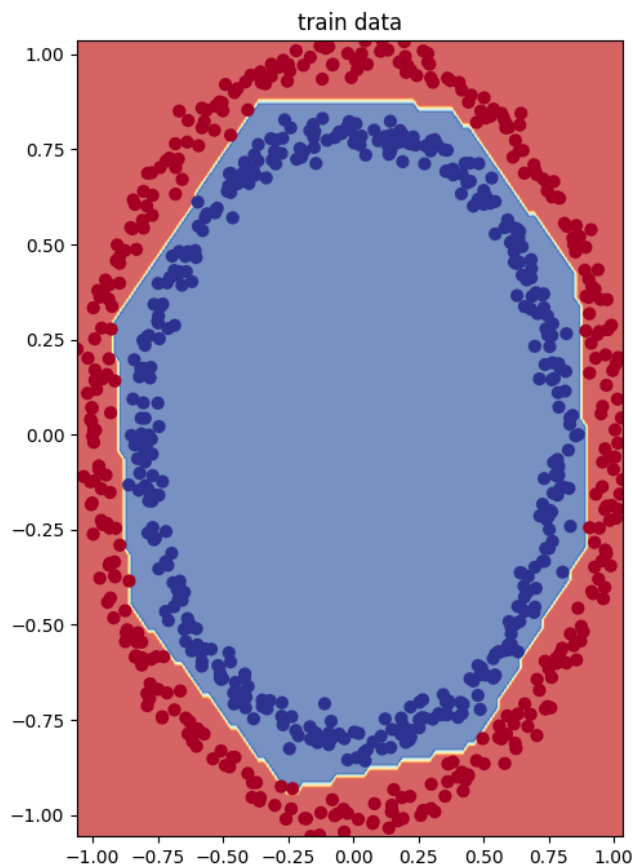
```
history=model_7.fit(X_train,Y_train,epochs=100,verbose=0)
```

```
model_7.evaluate(X_test,Y_test)
```

↗ 7/7 ————— 0s 2ms/step - accuracy: 0.9877 - loss: 0.0257
[0.028310153633356094, 0.9850000143051147]

```
plt.figure(figsize=(12,8))
plt.subplot(1,2,1)
plt.title("train data")
plot_decision_boundary(model_7,X_train,Y_train)
plt.subplot(1,2,2)
plt.title("test data")
plot_decision_boundary(model_7,X_test,Y_test)
```

↔ 313/313 ————— 1s 2ms/step
 doing binary classification
 313/313 ————— 1s 2ms/step
 doing binary classification




✓ plot loss or training curves

```
pd.DataFrame(history.history).plot()
```



<Axes: >



✓ Finding the best learning rate

use:

- a learning rate callback
- modified loss curve plot

```
tf.random.set_seed(42)
model_8=tf.keras.Sequential([
    tf.keras.layers.Dense(4,activation='relu'),
    tf.keras.layers.Dense(4,activation='relu'),
    tf.keras.layers.Dense(1,activation="sigmoid"),
])
model_8.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
                metrics=["accuracy"])

#create learning rate callback
lr_scheduler = tf.keras.callbacks.LearningRateScheduler(lambda epoch:1e-3*10**(epoch/20))

history_8=model_8.fit(X_train,Y_train,epochs=100,callbacks=[lr_scheduler])
```



```

Epoch 75/100
25/25 ————— 0s 3ms/step - accuracy: 0.5129 - loss: 0.7002 - learning rate: 0.001
Epoch 76/100
25/25 ————— 0s 3ms/step - accuracy: 0.5156 - loss: 0.7090 - learning rate: 0.001
Epoch 77/100
25/25 ————— 0s 3ms/step - accuracy: 0.5120 - loss: 0.7087 - learning rate: 0.001
Epoch 78/100
25/25 ————— 0s 3ms/step - accuracy: 0.5219 - loss: 0.7270 - learning rate: 0.001
Epoch 79/100
25/25 ————— 0s 3ms/step - accuracy: 0.5054 - loss: 0.7536 - learning rate: 0.001
Epoch 80/100
25/25 ————— 0s 3ms/step - accuracy: 0.5004 - loss: 0.7745 - learning rate: 0.001
Epoch 81/100
25/25 ————— 0s 3ms/step - accuracy: 0.4986 - loss: 0.7869 - learning rate: 0.001
Epoch 82/100
25/25 ————— 0s 3ms/step - accuracy: 0.4945 - loss: 0.8646 - learning rate: 0.001
Epoch 83/100
25/25 ————— 0s 4ms/step - accuracy: 0.5233 - loss: 0.8815 - learning rate: 0.001
Epoch 84/100
25/25 ————— 0s 3ms/step - accuracy: 0.4981 - loss: 1.2398 - learning rate: 0.001
Epoch 85/100
25/25 ————— 0s 1ms/step - accuracy: 0.4916 - loss: 0.8288 - learning rate: 0.001
Epoch 86/100
25/25 ————— 0s 1ms/step - accuracy: 0.4846 - loss: 0.9371 - learning rate: 0.001
Epoch 87/100
25/25 ————— 0s 2ms/step - accuracy: 0.4915 - loss: 0.9992 - learning rate: 0.001
Epoch 88/100
25/25 ————— 0s 2ms/step - accuracy: 0.4932 - loss: 1.4662 - learning rate: 0.001
Epoch 89/100
25/25 ————— 0s 2ms/step - accuracy: 0.5066 - loss: 1.1138 - learning rate: 0.001
Epoch 90/100
25/25 ————— 0s 2ms/step - accuracy: 0.5027 - loss: 1.2253 - learning rate: 0.001
Epoch 91/100
25/25 ————— 0s 2ms/step - accuracy: 0.5163 - loss: 2.2288 - learning rate: 0.001
Epoch 92/100
25/25 ————— 0s 2ms/step - accuracy: 0.4988 - loss: 2.0046 - learning rate: 0.001
Epoch 93/100
25/25 ————— 0s 2ms/step - accuracy: 0.4986 - loss: 2.1678 - learning rate: 0.001
Epoch 94/100
25/25 ————— 0s 1ms/step - accuracy: 0.4926 - loss: 2.4326 - learning rate: 0.001
Epoch 95/100
25/25 ————— 0s 1ms/step - accuracy: 0.5040 - loss: 7.2756 - learning rate: 0.001
Epoch 96/100
25/25 ————— 0s 2ms/step - accuracy: 0.5093 - loss: 10.8878 - learning rate: 0.001
Epoch 97/100
25/25 ————— 0s 2ms/step - accuracy: 0.5315 - loss: 10.6321 - learning rate: 0.001
Epoch 98/100
25/25 ————— 0s 2ms/step - accuracy: 0.5020 - loss: 11.9537 - learning rate: 0.001
Epoch 99/100
25/25 ————— 0s 2ms/step - accuracy: 0.4984 - loss: 3.3243 - learning rate: 0.001
Epoch 100/100
25/25 ————— 0s 2ms/step - accuracy: 0.4994 - loss: 3.4776 - learning rate: 0.001

```

```
model_8.evaluate(X_test,Y_test)
```

```

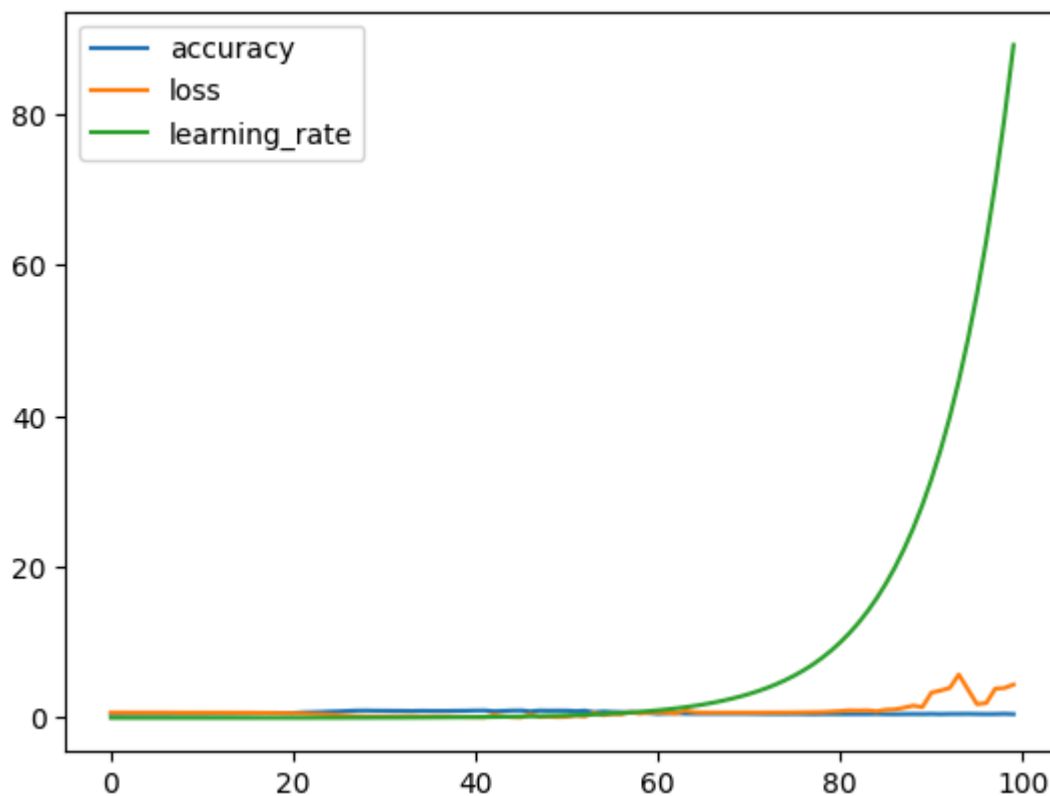
7/7 ————— 0s 2ms/step - accuracy: 0.4112 - loss: 5.0090
[4.934600830078125, 0.41999998688697815]

```

```
pd.DataFrame(history_9.history).plot()
```



<Axes: >



```
lrs = 1e-4 * (10 ** (tf.range(len(history_9.history['loss']) / 20)))
```

```
print("Shape of lrs:", lrs.shape)
print("Shape of loss:", len(history_9.history['loss']))
```



```
Shape of lrs: (5,)
Shape of loss: 100
```

Generate


a slider using jupyter widgets



Close

```
min_length = min(len(lrs), len(history_9.history['loss']))
lrs = lrs[:min_length]
loss = history_9.history['loss'][:min_length]
```

```
plt.figure(figsize=(10, 7))
plt.semilogx(lrs, loss)
plt.xlabel("Learning Rate (log scale)")
plt.ylabel("Loss")
plt.title("Learning Rate vs Loss")
plt.show()
```

- $\text{accuracy} = \frac{\text{tp} + \text{tn}}{\text{tp} + \text{tn} + \text{fp} + \text{fn}}$
- $\text{precision} = \frac{\text{tp}}{\text{tp} + \text{fp}}$ (high value leads to less fp)
- $\text{recall} = \frac{\text{tp}}{\text{tp} + \text{fn}}$ (high val leads to less fn)
- $\text{f1} = 2 * \text{precision} * \text{recall} / (\text{precision} + \text{recall})$
- confusion matrix
- classification report

there is a precision recall tradeoff

```
y_pred = model_7.predict(X_test)
```



7/7 ————— 0s 5ms/step

```
y_pred
```



```
array([[1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
        1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0,
        1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0,
        0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0,
        0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1,
        0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1,
        0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1])
```

```
0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0,  
1, 1])
```

our predictions has come out in **prediction probability** from the std output from sigmoid (or softmax)activation function

```
#convert prdiction probaability to binary from  
tf.round(y_pred)
```



```
[0.],
[1.],
[0.],
[1.],
[0.],
[0.],
[1.],
[0.],
[1.],
[1.]], dtype=float32)>
```

```
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, prec
```


```
confusion_matrix(Y_test,tf.round(y_pred))
```

```
↵ array([[114,  2],
        [ 1, 83]])
```

```
import seaborn as sns
```

```
sns.heatmap(confusion_matrix(Y_test,tf.round(y_pred)),annot=True)
```

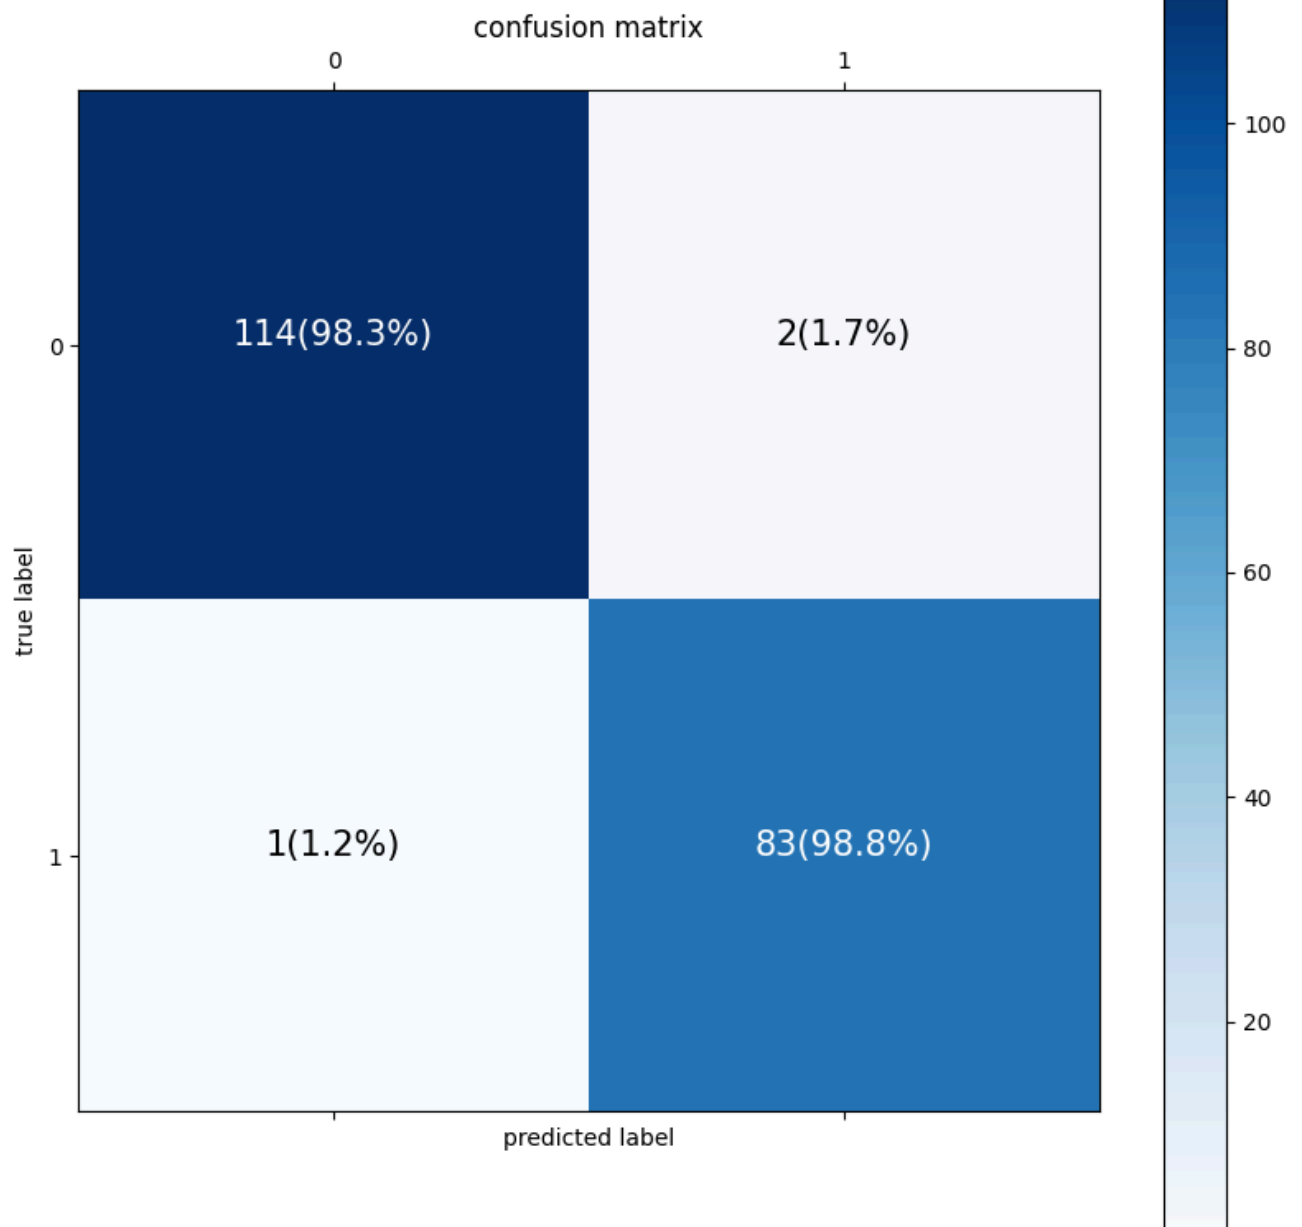
```
↵ <Axes: >
```



Suggested code may be subject to a licence | 1mampurwad1/mnist_cnn | CrispinGari/keras-api

```
import itertools
figsize=(10,10)
cm = confusion_matrix(Y_test,tf.round(y_pred))
cm_norm= cm.astype("float")/cm.sum(axis=1)[:,np.newaxis]
n_classes=cm.shape[0]
fig,ax=plt.subplots(figsize=figsize)
cax=ax.matshow(cm,cmap=plt.cm.Blues)
fig.colorbar(cax)
classes= False
```

```
if classes:
    labels=classes
else:
    labels=np.arange(cm.shape[0])
ax.set(title="confusion matrix",
        xlabel="predicted label",
        ylabel="true label",
        xticks=np.arange(n_classes),
        yticks=np.arange(n_classes),
        xticklabels=labels,
        yticklabels=labels)
threshold =(cm.max()+cm.min())/2
for i,j in itertools.product(range(cm.shape[0]),range(cm.shape[1])):
    plt.text(j,i,f"{cm[i,j]}({cm_norm[i,j]*100:.1f}%)",
            horizontalalignment="center",
            color="white" if cm[i,j]>threshold else "black",
            size=15)
```



✓ MULTICLASS CLASSIFICATION


- more than 2 different classes
- classify cloths is what we are gonna do

```
import tensorflow as tf
from tensorflow.keras.datasets import fashion_mnist
```

```
(train_data,train_labels),(test_data,test_lables)=fashion_mnist.load_data()
```

```
train_data[0]
```

↗ ndarray (28, 28) show data




```
train_data[0].shape,train_labels[0].shape
```

↗ ((28, 28), ())

```
import matplotlib.pyplot as plt
plt.imshow(train_data[5])
```

↗ <matplotlib.image.AxesImage at 0x7b817ca70d00>




- ✓ create a list so we can index onto our training labels so they are human readable

```
class_names=['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt', 'Sneaker'
```

```
plt.imshow(train_data[0], cmap=plt.cm.binary)
plt.title(class_names[train_labels[0]])
```

```
➞ Text(0.5, 1.0, 'Ankle boot')
```



```
plt.figure(figsize=(7,7))
for i in range(4):
    ax=plt.subplot(2,2,i+1)

    plt.imshow(train_data[i], cmap=plt.cm.binary)
    plt.title(class_names[train_labels[i]])
    plt.axis("off")
```



Ankle boot




T-shirt/top



T-shirt/top



Dress



✓ multiclass classification model

- input shape=28x28
- o/p shape=10
- loss func= categorical_crossentropy
- o/p layer activation=softmax

```
tf.random.set_seed(42)
model_10=tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(10,activation="softmax")
])
model_10.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=["accuracy"])
```


→ /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: Use super().__init__(**kwargs)

- CategoricalCrossentropy: one hot encoded values are expected
- otherwise use sparseCategoricalCrossEntropy(int type)

```
non_norm_history=model_10.fit(train_data,train_labels,epochs=10,validation_data=(test_data,
```

→ Epoch 1/10
 1875/1875 ————— 5s 2ms/step - accuracy: 0.1301 - loss: 3.1902 - val_acc: 0.1301
 Epoch 2/10
 1875/1875 ————— 4s 2ms/step - accuracy: 0.1995 - loss: 1.9566 - val_acc: 0.1995
 Epoch 3/10
 1875/1875 ————— 5s 3ms/step - accuracy: 0.3079 - loss: 1.6457 - val_acc: 0.3079
 Epoch 4/10
 1875/1875 ————— 9s 2ms/step - accuracy: 0.3898 - loss: 1.3959 - val_acc: 0.3898
 Epoch 5/10
 1875/1875 ————— 6s 2ms/step - accuracy: 0.4501 - loss: 1.2570 - val_acc: 0.4501
 Epoch 6/10
 1875/1875 ————— 4s 2ms/step - accuracy: 0.4741 - loss: 1.1832 - val_acc: 0.4741
 Epoch 7/10
 1875/1875 ————— 4s 2ms/step - accuracy: 0.4740 - loss: 1.1565 - val_acc: 0.4740
 Epoch 8/10
 1875/1875 ————— 5s 2ms/step - accuracy: 0.4869 - loss: 1.1422 - val_acc: 0.4869
 Epoch 9/10
 1875/1875 ————— 5s 2ms/step - accuracy: 0.4910 - loss: 1.1313 - val_acc: 0.4910
 Epoch 10/10
 1875/1875 ————— 5s 2ms/step - accuracy: 0.4954 - loss: 1.1243 - val_acc: 0.4954

🔗 Generate

create a dataframe with 2 columns and 10 rows



Close

```
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
scaled_train_data=scaler.fit_transform(train_data.reshape(-1,28*28))
scaled_test_data=scaler.transform(test_data.reshape(-1,28*28))


tf.random.set_seed(42)
model_11=tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(10,activation="softmax")
])
model_11.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=["accuracy"])
```

→ /keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`


```
train_data_norm=train_data/255.0  
test_data_norm=test_data/255.0
```

 **Generate****Close**

```
tf.random.set_seed(42)  
model_12=tf.keras.Sequential([  
    tf.keras.layers.Flatten(input_shape=(28,28)),  
    tf.keras.layers.Dense(4,activation="relu"),  
    tf.keras.layers.Dense(4,activation="relu"),  
    tf.keras.layers.Dense(10,activation="softmax")  
)  
model_12.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                 optimizer=tf.keras.optimizers.Adam(),  
                 metrics=["accuracy"])
```

 `/usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: Use
super().__init__(**kwargs)`

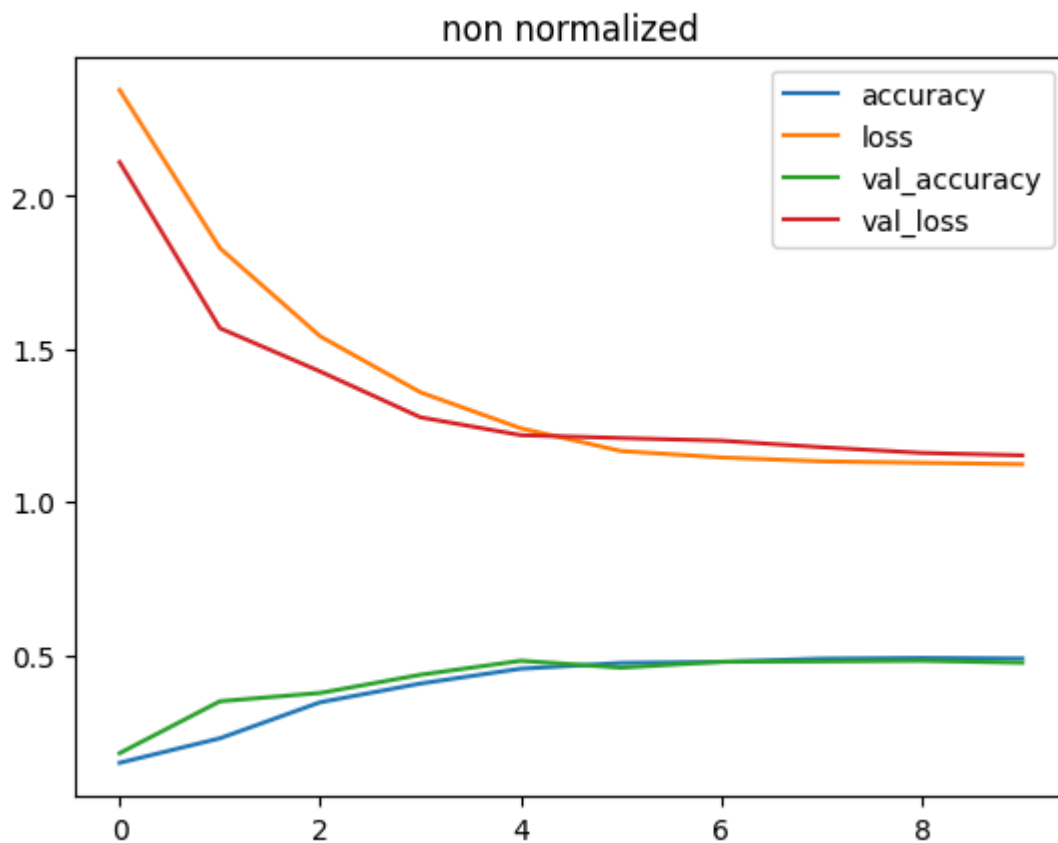

```
norm_history=model_12.fit(train_data_norm,train_labels,epochs=10,validation_data=(test_da
```



```
Epoch 1/10  
1875/1875 ————— 12s 5ms/step - accuracy: 0.3250 - loss: 1.8457 - val_ac  
Epoch 2/10  
1875/1875 ————— 4s 2ms/step - accuracy: 0.6243 - loss: 1.1099 - val_ac  
Epoch 3/10  
1875/1875 ————— 5s 2ms/step - accuracy: 0.6973 - loss: 0.8191 - val_ac  
Epoch 4/10  
1875/1875 ————— 5s 3ms/step - accuracy: 0.7233 - loss: 0.7226 - val_ac  
Epoch 5/10  
1875/1875 ————— 9s 2ms/step - accuracy: 0.7391 - loss: 0.6798 - val_ac  
Epoch 6/10  
1875/1875 ————— 5s 3ms/step - accuracy: 0.7525 - loss: 0.6525 - val_ac  
Epoch 7/10  
1875/1875 ————— 4s 2ms/step - accuracy: 0.7696 - loss: 0.6299 - val_ac  
Epoch 8/10  
1875/1875 ————— 5s 2ms/step - accuracy: 0.7781 - loss: 0.6139 - val_ac  
Epoch 9/10  
1875/1875 ————— 5s 3ms/step - accuracy: 0.7835 - loss: 0.6024 - val_ac  
Epoch 10/10  
1875/1875 ————— 4s 2ms/step - accuracy: 0.7879 - loss: 0.5938 - val_ac
```

```
import pandas as pd  
pd.DataFrame(norm_history.history).plot(title="normalized")  
pd.DataFrame(non_norm_history.history).plot(title="non normalized")
```

↔ <Axes: title={'center': 'non normalized'}>



compare with same criteria since small change has dramatic effects

```
#finding ideal learning rate
tf.random.set_seed(42)
model_13=tf.keras.Sequential([
```

```
tf.keras.layers.Flatten(input_shape=(28,28)),
tf.keras.layers.Dense(4,activation="relu"),
tf.keras.layers.Dense(4,activation="relu"),
tf.keras.layers.Dense(10,activation="softmax")
])
model_13.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=["accuracy"])
lr_scheduler=tf.keras.callbacks.LearningRateScheduler(lambda epoch:1e-3*10**(epoch/20))

↗ /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: Use
super().__init__(**kwargs)
```

```
model_13.fit(train_data_norm,train_labels,epochs=100,callbacks=[lr_scheduler])
```



```
Epoch 94/100
1875/1875 ————— 3s 2ms/step - accuracy: 0.0999 - loss: 19.8436 - le
Epoch 95/100
1875/1875 ————— 3s 2ms/step - accuracy: 0.0982 - loss: 24.5708 - le
Epoch 96/100
1875/1875 ————— 7s 2ms/step - accuracy: 0.0999 - loss: 27.7266 - le
Epoch 97/100
1875/1875 ————— 4s 2ms/step - accuracy: 0.1009 - loss: 31.2183 - le
Epoch 98/100
1875/1875 ————— 3s 2ms/step - accuracy: 0.0993 - loss: 36.5593 - le
Epoch 99/100
1875/1875 ————— 6s 2ms/step - accuracy: 0.1016 - loss: 42.5347 - le
Epoch 100/100
1875/1875 ————— 4s 2ms/step - accuracy: 0.1011 - loss: 44.4787 - le
<keras.src.callbacks.history.History at 0x7b8177b7af20>
```


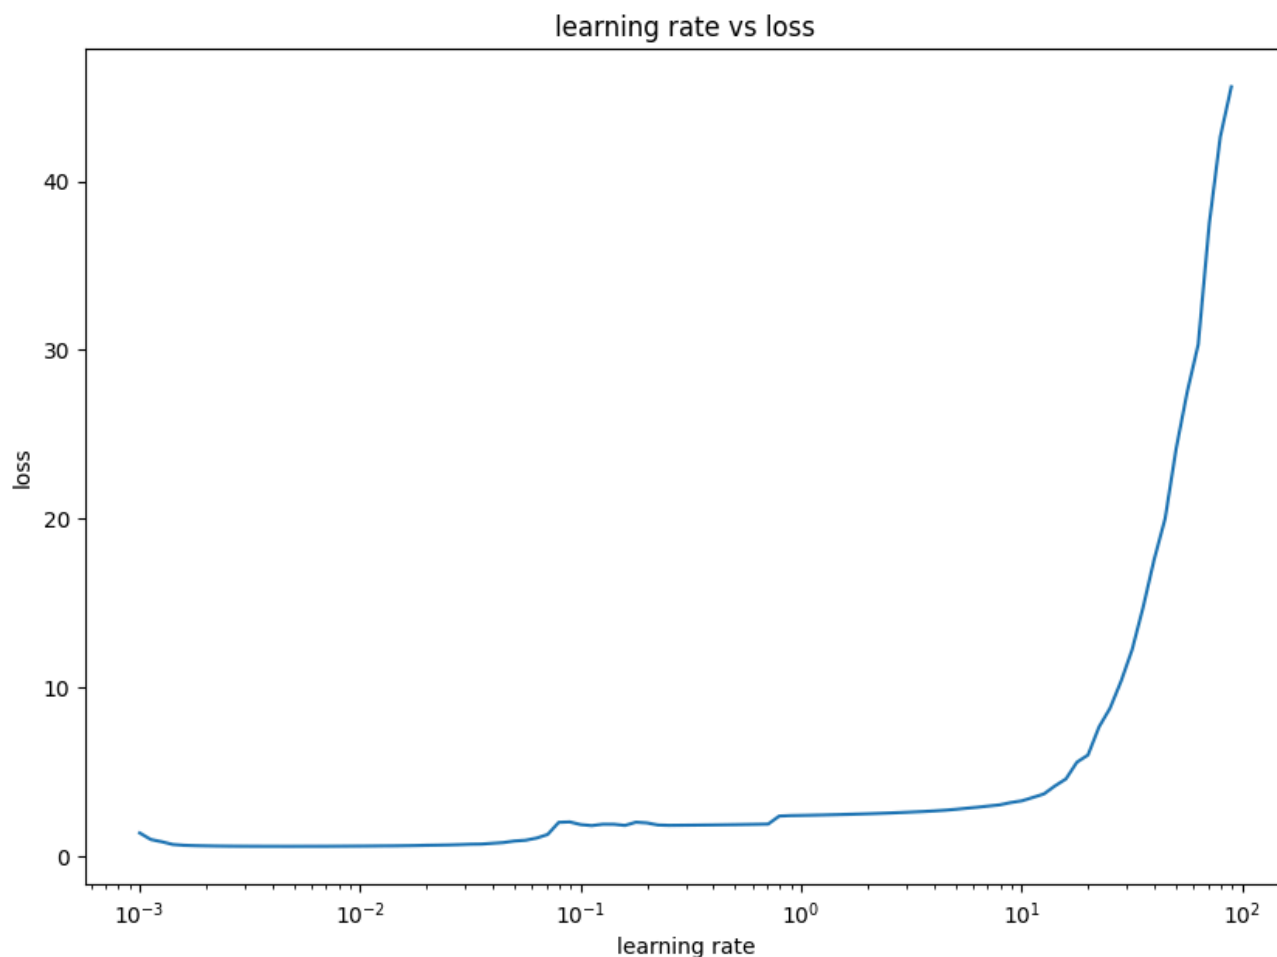
 Generate

a slider using jupyter widgets



Close

```
import numpy as np
import matplotlib.pyplot as plt
lrs=1e-3*(10**(np.arange(100)/20))
plt.figure(figsize=(10,7))
plt.semilogx(lrs,model_13.history.history["loss"])
plt.xlabel("learning rate")
plt.ylabel("loss")
plt.title("learning rate vs loss")
```












 Text(0.5, 1.0, 'learning rate vs loss')

ideal learning rate is 0.001

```
tf.random.set_seed(42)
model_14=tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(10,activation="softmax")
])
model_14.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                 optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
                 metrics=["accuracy"])
```

 /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: Use
super().__init__(**kwargs)

history_14 = model_14.fit(train_data_norm,train_labels,epochs=10,validation_data=(test_da

 Epoch 1/10
1875/1875  6s 2ms/step - accuracy: 0.4281 - loss: 1.5809 - val_ac
Epoch 2/10
1875/1875  4s 2ms/step - accuracy: 0.7442 - loss: 0.7574 - val_ac
Epoch 3/10
1875/1875  6s 2ms/step - accuracy: 0.7778 - loss: 0.6540 - val_ac
Epoch 4/10
1875/1875  4s 2ms/step - accuracy: 0.7902 - loss: 0.6050 - val_ac
Epoch 5/10
1875/1875  4s 2ms/step - accuracy: 0.8004 - loss: 0.5767 - val_ac
Epoch 6/10
1875/1875  5s 2ms/step - accuracy: 0.8066 - loss: 0.5579 - val_ac
Epoch 7/10
1875/1875  6s 2ms/step - accuracy: 0.8115 - loss: 0.5444 - val_ac
Epoch 8/10
1875/1875  5s 2ms/step - accuracy: 0.8147 - loss: 0.5338 - val_ac
Epoch 9/10
1875/1875  7s 3ms/step - accuracy: 0.8168 - loss: 0.5263 - val_ac
Epoch 10/10
1875/1875  4s 2ms/step - accuracy: 0.8192 - loss: 0.5202 - val_ac

✓ EVALUATING OUR MODEL WITH OTHER METRICS

- classification metrics
- asses some of its predictions
- improve its results(change architechture)
- save and export

```
import itertools
from sklearn.metrics import confusion_matrix
```

```
def make_conf_matrix(y_true,y_pred,classes= None,figsize=(10,10)):
    cm = confusion_matrix(y_true,y_pred)
    cm_norm= cm.astype("float")/cm.sum(axis=1)[: ,np.newaxis]
    n_classes=cm.shape[0]
    fig,ax=plt.subplots(figsize=figsize)
    cax=ax.matshow(cm,cmap=plt.cm.Blues)
    fig.colorbar(cax)


    if classes:
        labels=classes
    else:
        labels=np.arange(cm.shape[0])
    ax.set(title="confusion matrix",
          xlabel="predicted label",
          ylabel="true label",
          xticks=np.arange(n_classes),
          yticks=np.arange(n_classes),
          xticklabels=labels,
          yticklabels=labels)
    threshold =(cm.max()+cm.min())/2
    for i,j in itertools.product(range(cm.shape[0]),range(cm.shape[1])):
        plt.text(j,i,f"{cm[i,j]}({cm_norm[i,j]*100:.1f}%)",
                horizontalalignment="center",
                color="white" if cm[i,j]>threshold else "black",
                size=15)
```

```
y_probs=model_14.predict(test_data_norm)
```

 313/313 ————— 0s 1ms/step

NOTE: remember to make predictions on the same kind of data your model was trained on (eg if trained on normalized data use normalized data for predictions)

```
y_probs[0],tf.argmax(y_probs[0])
```

 (array([1.5085043e-03, 1.1936341e-04, 1.0228357e-03, 8.6857857e-05,
1.2778917e-05, 3.5110056e-01, 2.8528241e-04, 6.0521282e-02,
3.2882905e-05, 5.8530962e-01], dtype=float32),
<tf.Tensor: shape=(), dtype=int64, numpy=9>)

```
k=tf.argmax(y_probs[0]).numpy()
```

 **Generate**

a slider using jupyter widgets



Close

```
class_names[k]
```

 'Ankle boot'

```
#convert all prediction pobabilities to integers
y_preds=y_probs.argmax(axis=1)
y_preds
```

```
array([9, 2, 1, ..., 4, 1, 5])
```

 **Generate**

a slider using jupyter widgets




Close

```
from sklearn.metrics import confusion_matrix
confusion_matrix(test_labels,y_preds)
```

```
array([[778,  4, 19, 58,  1,  2, 113,  2, 20,  3],
       [ 16, 923,  3, 47,  1,  0,  6,  2,  2,  0],
       [ 27,  1, 794, 14, 61,  0, 96,  0,  7,  0],
       [ 30, 17, 18, 852, 22,  3, 45,  1, 12,  0],
       [  1,  0, 235, 57, 553,  0, 143,  0, 11,  0],
       [  0,  0,  1,  0,  0, 902,  0, 52,  3, 42],
       [183,  0, 180, 44, 67,  0, 507,  0, 19,  0],
       [  0,  0,  0,  0,  0, 44,  0, 915,  0, 41],
       [  4,  2,  8,  9,  8, 24, 27,  4, 914,  0],
       [  0,  0,  0,  0,  0, 20,  0, 45,  1, 934]])
```

```
make_conf_matrix(test_labels,y_preds,classes=class_names)
```





```
import random
def plot_random_image(model,images,true_labels,classes):
    """
    Picks a random image, plots it and labels it with a prediction and truth label
    """
    i=random.randint(0,len(images))
    target_image=images[i]
    pred_probs=model.predict(target_image.reshape(1,28,28))
    pred_label=class_names[pred_probs.argmax()]
    true_label=class_names[true_labels[i]]
    plt.imshow(target_image,cmap=plt.cm.binary)
    if pred_label == true_label:
        color="green"
    else:
        color="red"
    plt.xlabel(f"pred: {pred_label} \n true: {true_label}",color=color)
    plt.title(f"pred: {pred_label} \n true: {true_label}")
    plt.axis("off")
```

```
plot_random_image(model_14, test_data_norm, test_labels, class_names)
```

1/1 0s 28ms/step

pred: Bag
true: Bag



```
model_14.layers[1]
```

<Dense name=dense_53, built=True>

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.